Open-Weight Instruction-Tuned and Edge-Capable LLMs for Greek NLP: A Reproducible Evaluation

Anonymous Author(s)

ABSTRACT

Natural Language Processing (NLP) for lesser-resourced languages faces persistent challenges, including fragmented datasets, inherited biases from high-resource settings, and limited reproducible evaluation protocols. This study addresses these gaps for Modern Greek through three contributions. First, we distill publicly available Greek datasets from prior survey work and classify them according to their suitability for benchmarking, pre-training, and retrieval-augmented generation. Second, we evaluate open instruction-tuned large language models (LLMs) on seven Greek NLP task families for which data are available and legally reusable, revealing task-specific strengths and weaknesses across classification, sequence labeling, translation, correction, and summarization. Third, we provide a reproducible benchmark pipeline that connects dataset selection, model execution, metric reporting, publication-ready artifacts, and automated discovery of candidate datasets from newer literature. Together, these contributions support more transparent evaluation and reuse of Greek NLP resources.

CCS CONCEPTS

• **Computing methodologies** → **Natural language processing**; *Language resources*; • **Information systems** → Data analytics.

KEYWORDS

Greek NLP, benchmarking, large language models, reproducibility, evaluation

ACM Reference Format:
Anonymous Author(s). 2026. Open-Weight Instruction-Tuned and Edge-Capable LLMs for Greek NLP: A Reproducible Evaluation. In *Proceedings of 14th Hellenic Conference on Artificial Intelligence (SETN '26).* ACM, New York, NY, USA, 15 pages. https://doi.org/XXXXXXX.XXXXXXX

1 INTRODUCTION

Natural Language Processing (NLP) has advanced substantially with the rise of large language models (LLMs), whose development depends on immense volumes of digital data [5]. Although multilingual NLP has benefited from these advances, progress for lesser-resourced languages continues to lag behind that of well-supported languages. As a result, NLP for many of the world's languages often relies on methods, resources, and evaluation practices developed for high-resource settings, inheriting assumptions, biases, and other characteristics that may not align with the properties of less supported languages [**?**].

This gap is compounded by the scarcity of open-access, high-quality language resources. In this context, high-quality resources refer to native, human-authored corpora, as opposed to synthetic or machine-translated data. Although synthetic and translated resources are often used to bridge data gaps, they can introduce "translationese", where the text inherits syntactic and stylistic constraints from the source language [17], and may lack the lexical diversity and morphological nuance of original language use. Such resources are needed both for pre-training foundation models and for building gold-standard benchmarks for downstream tasks. Comprehensive surveys of NLP resources for lesser-resourced languages therefore do more than document digital readiness: they can also guide future research, reveal reusable datasets, and support the construction of new benchmarks. Our work focuses on one such language, Modern Greek.[1]

In this work, we ask two related research questions: *which properly licensed Greek datasets can be used to train and/or assess LLMs on downstream NLP tasks, and how do open, instruction-tuned, edge-capable LLMs perform across the resulting Greek benchmark?* By addressing these questions, our work makes three contributions:

- Based on a recent Greek NLP survey [**?**], we extract publicly available Greek datasets that are (a) open-access and (b) permit derivatives, aggregate them into a unified collection that is easy to recompile and use, and provide an automated discovery component for identifying candidate datasets from newer Greek NLP literature.[2]
- We classify the collected datasets according to their usability for benchmarking or pre-training.
- We benchmark open LLMs with manageable parameter counts and Greek-language capabilities across multiple NLP tasks, revealing task-specific strengths and weaknesses.

2 RELATED WORK

2.1 Greek NLP Resources and Models

Modern Greek NLP has produced a sustained body of task-specific resources and systems over the past decades, including work on grammar checking, legal text processing, named entity recognition, sentiment analysis, and syntactic annotation [1, 3, 11, 26, 31]. More recent surveys have shown, however, that these resources remain unevenly distributed across tasks and domains, and that availability, licensing, and reusability are recurring barriers for reproducible experimentation [24**?**]. This motivates our focus on datasets that are not only relevant to downstream Greek NLP, but also accessible and legally reusable.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
SETN '26, September 9–11, 2026, Chania, Greece
© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM
https://doi.org/XXXXXXX.XXXXXXX

[1] In our study, the term "Greek" refers exclusively to Modern Greek (ISO 639-1: el), excluding previous stages of the language's history.
[2] The code is available at: anonymised.url

Transformer-based models have also expanded the practical toolkit for Greek. Greek-specific pretrained models such as GreekBERT and GreekBART improved representation learning and sequence-to-sequence modelling for the language [9, 19], while recent open LLMs for Greek, such as Meltemi, indicate growing interest in language-specialised generative models [38]. In parallel, multilingual and open instruction-tuned LLMs such as Llama, Gemma, Aya, Qwen, and Falcon provide broader coverage and are increasingly used through local serving frameworks [2, 7, 14, 36, 42]. Our work complements these model-development efforts by evaluating manageable open models under a shared Greek benchmark protocol.

2.2 Benchmarking and Evaluation

Multi-task benchmarks such as GLUE and SuperGLUE established shared evaluation as a central mechanism for comparing general-purpose language understanding systems [39, 40]. Similar efforts have since appeared for individual languages, domains, and multilingual settings, including legal and lower-resource evaluation suites [6, 37, 41]. For foundation models and LLMs, broad evaluation frameworks such as HELM and the LM Evaluation Harness further emphasize reproducibility, task coverage, and transparent reporting [12, 20].

Despite this progress, Greek LLM evaluation remains fragmented across datasets, task definitions, metrics, and execution environments. Existing Greek studies often target individual tasks or model families, making it difficult to compare open instruction-tuned LLMs across heterogeneous downstream settings. Our benchmark therefore prioritizes a reproducible evaluation pipeline over a new modelling method: it brings together properly licensed Greek datasets, standardizes task execution and reporting, and produces outputs that can be inspected, reproduced, and reused in subsequent work.

3 THE DATA: FROM COLLECTION TO DISTILLATION

Based on a systematic literature review covering 142 research NLP studies for Greek from 2012 to 2023 [**?**], we extracted and analysed information on 94 publicly available Greek datasets. Most of these resources were monolingual (84%), but bilingual and multilingual datasets were also present (i.e., with Greek being one of the languages). The most frequent domains were politics and news (20%) while the most frequent task was sentiment analysis (33%).

3.1 Availability, Popularity and Licensing

Figure 1 shows the availability status of the collected datasets. The most common status is that **the datasets are not available**, meaning no information is provided by the authors of the respective study regarding their availability. The second most frequent category concerns datasets reported as downloadable, but **actually not publicly available** due to restrictions, such as missing license or subscription limitations. This limitation contradicts the principles and benefits of open access.

Studies with publicly available datasets, serving as a reference point, tend to attract more citations. This is shown by the higher third and fourth quartile of the green box in Figure 2. Furthermore, studies attracting citations while presenting data with limited [15] or no availability [13] belong to the past. The study of Juola and Stamatatos [15] is such an example, shown in Figure 3 with a blue peak in 2013, concerning a shared task on authorship identification (hence, the citations) whose data lack a license (hence, the limitation).

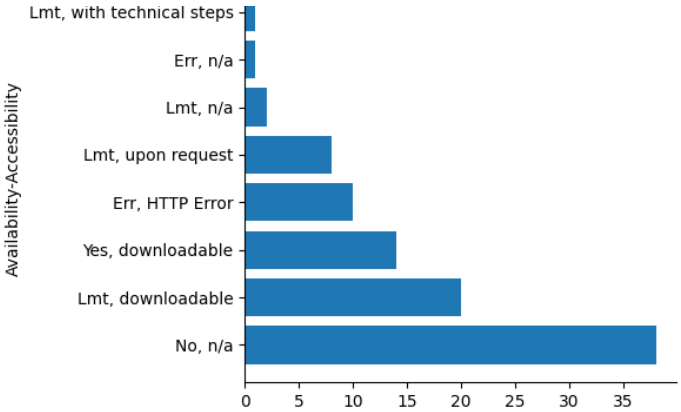

Figure 1: Availability and accessibility of collected datasets (tuple per bar). Availability is classified to Yes, Lmt (no license or otherwise constrained), Err (hindered by an error), No (not available via a URL). Accessibility is classified to downloadable, HTTP error, upon request, n/a (missing resource); i.e., the outcome we observed when accessing the resource.

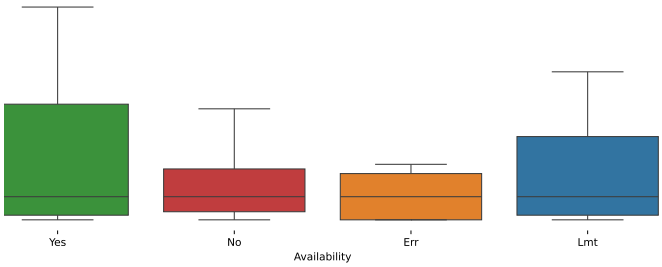

Figure 2: Whisker and box plots with the number of citations of studies per availability type.

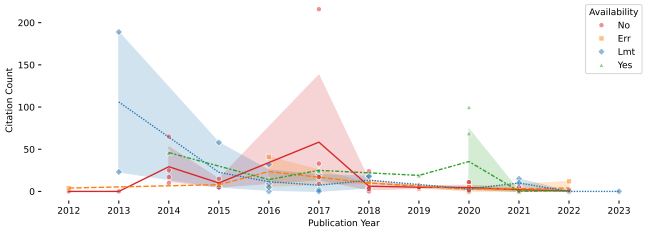

Figure 3: Citation counts per year of publication of studies developing datasets, based on the availability classified as yes, no, limited (lmt), or erroneous (err).

3.2 AIR Distillation

We collected accessible, interoperable and re-usable (AIR)[3] datasets pooled from a recent survey [?]. We focused on the 14 Greek datasets that meet our criteria of being licensed, accessible, machine-actionable, and with no hidden costs. We excluded one [10] that was licensed under CC BY-NC-ND 4.0, which does not allow derivatives and thus limits reusability. We also excluded one [43] that shared tweets not via rehydration (i.e., only IDs shared), but as full texts.[4] From the two resources released by Korre et al. [18], we retained the 227-sentence Greek Native Corpus (GNC) for GEC, but excluded the 100-sentence Greek WikiEdits corpus (GWE), which is too small for our study. We also excluded the dataset of open-access e-books [4], because it comprises OCR artifacts.[5] Given these exclusions, we end up with ten datasets, presented in Table 1, which cover a variety of domains.

3.3 AIR Resources Suitable for RAG

Table 1 comprises three datasets with no ground truth or whose ground truth is automatically extracted (A/N; i.e., of uncertain quality). Although excluded from our benchmark, these datasets are still possibly useful for retrieval-augmented generation (RAG) and other purposes (e.g., pre-training). These concern 1.28M political speeches, spanning from July 1989 to July 2020 [8]; Greek Wikipedia event pages with named entity recognition (NER) and linking; 101,857 (cleaned) websites with open content [33], including online archives of Greek newspapers from 2003 to 2020, and the Greek part of the W2C corpus [22].

3.3.1 Dataset statistics. As shown in Table 2, the text length (measured in characters) greatly varies between the datasets. The dataset with the shortest texts on average is that of Papantoniou et al. [27] and that with the lengthiest is that of Prokopidis and Piperidis [33]. The dataset size (counted in number of files) also varies, starting from ten thousands [27], to one hundred thousands [33], to more than a million [8]. The aggregated text totals up to 1.4 billion characters and 211 million tokens (using white space split).

3.3.2 Linguistic distance. We investigate the linguistic associations between these datasets, using statistical language modelling (LMing) as a proxy. For each dataset, we trained one statistical character-level LM on the first 100 characters of 1,000 randomly selected texts per dataset.[6] Then, we computed the Bits per Characters (BPC) across 500 randomly selected texts per dataset, reporting the mean. The result is displayed as a heatmap in Figure 4. High scores (in red) reflecting a greater linguistic distance between the dataset the LM is trained on and the one BPC is computed on. Low scores (in blue) reflect a low linguistic distance, with the lowest scores taking place in the diagonal, when predicting unseen text taken from the same dataset (self-distance). Based on this heatmap, we find that the dataset of Prokopidis and Piperidis [33] shows a higher self-distance compared to the rest while the datasets of Dritsa et al. [8] and Papantoniou et al. [27] are far away from each other, especially when the LM is trained on the former and predicting the latter.

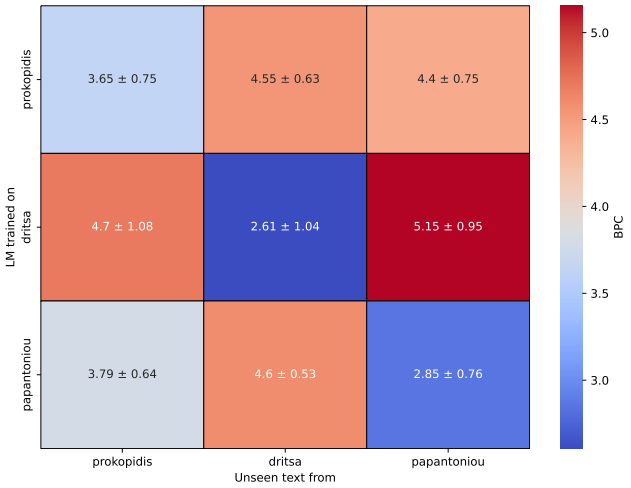

Figure 4: Heatmap of BPC measured on a sample per RAG dataset (average across the sample; shown horizontally) per trained LM (vertically); greater distance in warm colours.

3.4 Resources Suitable for Downstream Tasks

The datasets used as benchmarks for downstream Greek NLP tasks are summarized in Table 3, along with key statistics and their associated supervised tasks. These resources cover three main categories: classification (Intent, Legal), sequence labeling (POS tagging, NER), and text-to-text generation (Grammatical Error Correction, Machine Translation, Summarization). Next, we present each one.

3.4.1 Intent Classification. Student queries, addressed to two university help desks, were manually annotated with six intent labels and three entity types [34]. This dataset is balanced across intent classes, with relatively short texts (average length: 58.6 characters). We use it to evaluate intent classification.

3.4.2 Legal Text Classification. A multi-level legal classification dataset is derived from the "Permanent Greek Legislation Code – Raptarchis" [26]. It includes hierarchical annotations of 47 volumes, 347 chapters, and 1,685 subjects. The dataset is highly imbalanced and exhibits large variation in document length (from 17 to approx. 1 million characters). The largest class comprises 34 examples (6.8%; science and arts) and the median has 10 examples. The smallest class has two examples (0.4%; media and tourism). We use it for volume classification.

3.4.3 POS Tagging. The Greek Universal Dependencies treebank was part of a project offering standardized treebanks with consistent annotations across languages [25]. This dataset [31] includes

[3] To comply with the FAIR principles, datasets must have a unique and persistent digital identifier (commonly a Persistent Identifier/PID). A PID is a specific combination of numbers and letters that ensures datasets can be uniquely and sustainably accessed as digital objects over time. Merely providing a URL for accessibility does not make a dataset findable, as URLs can become obsolete or break if not properly maintained.

[4] Although it is possible that a company gained access to the texts (e.g., via firehose), releasing them with another license, a user should still be able to delete their content from all platforms. If not, then this dataset is not compliant with the rules. If so, then the dataset is not stable, harming reproducibility of experiments.

[5] Although typos and errors are not prohibited for experiments, we observe that the author names and the title was often injected in the main text, leading to data leaking for potential benchmarks.

[6] Character-level statistical LMing is an efficient way to capture stylistic features, employed in authorship analysis [28].

Table 1: Greek datasets with their annotation type (AT) defined as manual (M), automatic (A), curated (C; metadata provided by the distributor), hybrid (H; manual and automatic), user-generated (U; from user-edits, not curated) or no (N). We include also the host, size, size unit, domain and license of each dataset.

AUTHORS	AT	HOST	SIZE	SIZE UNIT	DOMAIN	LICENSE
Papantoniou et al. [27]	A	Zenodo	18,615	DOC	EVENTS	CC BY 3.0
Koniaris et al. [16]	*C*	Hugging Face	8,395	DOC	LEGAL	CC BY-NC 4.0
Rizou et al. [34]	*M*	msensis.com	2,176	SENT	ADMIN	CC-BY-NC-SA 4.0
Dritsa et al. [8]	A	Zenodo	1.28M	DOC	POLITIC	CC BY 4.0
Papaloukas et al. [26]	*C*	Hugging Face	47,563	DOC	LEGAL	CC BY 4.0
Korre et al. [18]	*M*	GitHub	227	SENT	ESSAYS	MIT
Bartziokas et al. [3]	*H*	GitHub	21,153	SENT	NEWS	CC-BY-NC-SA 4.0
Prokopidis and Piperidis [33]	N	nlp.ilsp.gr	101,857	DOC	WEB	CC-BY-NC-SA 4.0
Prokopidis and Papageorgiou [31]	*H*	GitHub	2,521	SENT	NEWS, POLITIC	CC BY-NC-SA 3.0
Prokopidis et al. [32]	*U*	nlp.ilsp.gr	17,018	DOC PAIR	NEWS	CC BY 4.0

Table 2: AIR datasets for RAG (not in the benchmark) with min, max, mean length in characters and size in documents.

	MIN	MAX	AVG	SIZE (#)
Papantoniou et al. [27]	34	1,484	172.9	18,615
Prokopidis and Piperidis [33]	36	67,963	2,430	101,857
Dritsa et al. [8]	1	291,949	900	1,280,927

syntactic dependencies, POS tags, morphological features, and lemmas. It consists of 2,521 manually validated sentences split into train (1,622), development (403), and test (456) sets.

3.4.4 NER. The Greek NER dataset [3] comprises two annotation schemes: a coarse-grained set of 4 entity types (similar to CONLL-2003 [35]) and a fine-grained set of 18 types (similar to OntoNotes [30]). The dataset was developed during a Google Summer of Code project in 2018,[7], where initial automatic annotation was followed by manual curation. The text length varies from 2 to 2,210 characters (Table 3). Both annotation levels exhibit heavy imbalance, with the four-tag level having an imbalance ratio (minor:major) of 0.02 and the 18-tag level showing near-zero imbalance. We assess NER at the less imbalanced four-entity-tag level.

3.4.5 Grammatical Error Correction (GEC). The Greek Native Corpus (GNC) [18] comprises (227) student-written sentences, annotated by human experts for (18) error types, related to grammar, spelling, punctuation, and morphology. Each sentence may contain zero, one, or multiple grammatical errors. We use this dataset to assess grammatical error correction.

3.4.6 Machine Translation (MT). Bilingual corpora (756 language pairs) exist from content available in Global Voices, where volunteers translate news stories in 41 languages [32]. The 3,629 Greek documents are translated into 40 languages. However, not every document is translated into all languages, resulting in 17,018 bilingual document pairs involving Greek. We use the translations to

[7]https://github.com/eellak/gsoc2018-spacy

Table 3: Dataset statistics for downstream Greek NLP tasks. We report test set size, number of classes (#C), average text length (in characters), and class imbalance ratio (CIR; smallest-to-largest class proportion) where applicable. Datasets split by us are marked with †.

TASK	TEST	UNIT	#CL.	AVG CHAR.	CIR
Intent Cl. [34]	436†	sent	6	58.6	1.00
Legal Text Cl. [26]	500	text	47	3673.3	0.07
POS Tagging [31]	456	sent	16	138.2	0.03
NER [3]	500†	sent	4	180.9	0.02
GEC [18]	175	sent	-	483.0	-
MT [32]	1,500†	doc	3	125.1	-
Summarization [16]	300	text	-	18,771.3	-

English, Japanese, and Farsi, representing high-, moderately-, and low-supported language tiers [**?**].

3.4.7 Summarization. A legal corpus of 8,395 court decisions is compiled from the Greek Supreme Court (Areios Pagos),[8] including the full texts, their summaries, case categories, and associated tags [16]. This dataset exhibits substantial length variability, with 18,541 characters on average and 18,351 st. dev. (min: 2,571; max: 303,538). We use it to assess summarization.

4 BENCHMARK SETUP

4.1 Tasks and Models

Our benchmark covers seven task families: GEC, intent classification, legal classification, machine translation, NER, POS tagging, and summarization.

We evaluate five open-weight, instruction-tuned models in the 7B–9B parameter range that are available for local inference through Ollama:

[8]https://www.areiospagos.gr/

- `qwen2.5:7b-instruct` [42];[9]
- `llama3.1:8b-instruct` [14];[10]
- `aya-expanse:8b` [7];[11]
- `gemma2:9b` [36];[12]
- `falcon3:7b-instruct` [2];[13]

This set controls for model scale while maximizing architectural and training diversity. Llama 3.1 and Qwen2.5 provide long-context variants (128K), Falcon3 reports 32K context, and Aya Expanse reports 8K context.

All models are evaluated under the same benchmark pipeline, prompts, dataset splits, and metrics, so observed differences mainly reflect model capability, multilingual transfer, and instruction-following quality rather than evaluation setup.

4.2 Deterministic Capped-Test Protocol

The full test sets are very uneven in size, with legal classification and machine translation being especially large. To keep evaluation tractable while preserving reproducibility, we use a deterministic capped-test profile that always takes the first instances in dataset order (as suggested by the dataset providers) instead of random samples. The resulting capped evaluation (Table 3) contains:

- GEC: 175 examples
- intent classification: 436 examples
- legal classification: 500 examples
- machine translation: 500 English pairs, 500 Persian pairs, and 500 Japanese pairs
- NER: 500 examples
- POS tagging: 456 examples
- summarization: 300 examples

This design keeps the smaller tasks intact while reducing the runtime burden of the largest tasks. It also supports exact reruns because the selected subset is deterministic. For machine translation, the benchmark now uses a shared data-wrapper path rather than a separate task-local split implementation, which eliminates an earlier source of inconsistency in language-specific counts.

4.3 Metrics

We report task-appropriate primary metrics. For GEC, the primary metric is GLEU [23]; we also report exact match and edit-distance diagnostics (WER/CER) against reference and source text. For intent classification, legal classification, NER, and POS tagging, macro-F1 is the primary comparison metric, with accuracy as a supplementary metric; for NER we additionally report entity-level F1. For machine translation we use chrF [29] as the primary metric and retain WER/CER as supplementary string-distance diagnostics. For summarization we use BERTScore-F1 [44] as the primary semantic comparison metric and report WER/CER and ROUGE [21] as supplementary overlap-based diagnostics. Latency is reported as average per-instance generation time in seconds.

[9]https://huggingface.co/Qwen/Qwen2.5-7B-Instruct
[10]https://huggingface.co/meta-llama/Llama-3.1-8B-Instruct
[11]https://huggingface.co/CohereLabs/aya-expanse-8b
[12]https://huggingface.co/google/gemma-2-9b
[13]https://huggingface.co/tiiuae/Falcon3-7B-Instruct

To compare models across heterogeneous metrics, we additionally compute a direction-aware normalized quality score per benchmark segment. Within each segment, the best model receives the highest normalized value and the worst model the lowest. These segment-normalized scores are then averaged to obtain an overall ranking. The main paper focuses on the primary metrics used for ranking, while the appendix reports additional prompt and winner-margin material.

5 EXPERIMENTAL RESULTS

We analyse the experimental results by focusing on the best model per task, but also touching on the relationship between performance and efficiency. Table 4 summarizes the best model in each benchmark segment. Figure 5 visualizes normalized quality across segments, while Figure 6 shows the trade-off between quality and inference speed. Furthermore, we provide an analytical table per task.

5.1 Patterns across models

5.1.1 The best model. The strongest overall model is `gemma2:9b`, ranking first on the normalized aggregate with an average quality score of 84.4 and winning intent classification, legal text classification, NER, POS tagging, and two of the three machine translation segments.

5.1.2 The runner up complements the best. `llama3.1:8b-instruct` ranks second overall with an average normalized quality score of 53.8. It wins GEC and remains highly competitive in intent classification and summarization, but it trails `gemma2:9b` on legal text classification and the aggregate ranking.

5.1.3 Competitive in cross-lingual tasks. `qwen2.5:7b-instruct` is third, which is especially competitive on Persian machine translation and wins that segment. `aya-expanse:8b` ranks fourth overall, but remains a strong specialist in summarization and is the runner-up on legal text classification and English translation.

5.1.4 The weakest model is not uniformly weak. The weakest model is `falcon3:7b-instruct`. This is often among the fastest models even when it lags in quality, which reinforces the value of reporting latency alongside task scores.

5.2 Patterns across tasks

5.2.1 Text classification. Legal text classification remains a difficult benchmark component. As shown in Table 4 and the detailed results in Table 7, the refreshed coarse-grained setup raises the best macro-F1 to 0.126 with `gemma2:9b`.[14] This suggests that the task is still demanding under zero-shot prompting and may benefit from better label descriptions, exemplars, or task-specific fine-tuning.

By contrast, intent classification is comparatively more tractable. As reported in Table 4 and Table 6, `gemma2:9b` achieves a macro-F1 score of 0.737, substantially outperforming other models on this task.

5.2.2 Translation. The capped profile evaluates 500 examples per target language, which is enough to separate models consistently while keeping runtime manageable. As summarized in Table 4 and

[14]We replaced opaque identifiers with coarse-grained volume labels.

Table 4: Best-performing model per benchmark segment on the deterministic capped evaluation run (GEC=175, Intent=436, Legal=500, MT(ENG=500, FAS=500, JPN=500), NER=500, POS=456, Summarization=300), using representative metrics per task.

Task Segment	Winner	Primary Metric	Value
GEC	Llama 3.1 8B Instruct	GLEU	0.796
Intent	Gemma2 9B	Macro-F1	0.737
Legal	Gemma2 9B	Macro-F1	0.126
MT-ENG	Gemma2 9B	chrF	61.781
MT-FAS	Qwen2.5 7B Instruct	chrF	18.325
MT-JPN	Gemma2 9B	chrF	16.726
NER	Gemma2 9B	Macro-F1	0.122
POS	Gemma2 9B	Macro-F1	0.486
Summ.	Aya Expanse 8B	BERTScore-F1	0.519

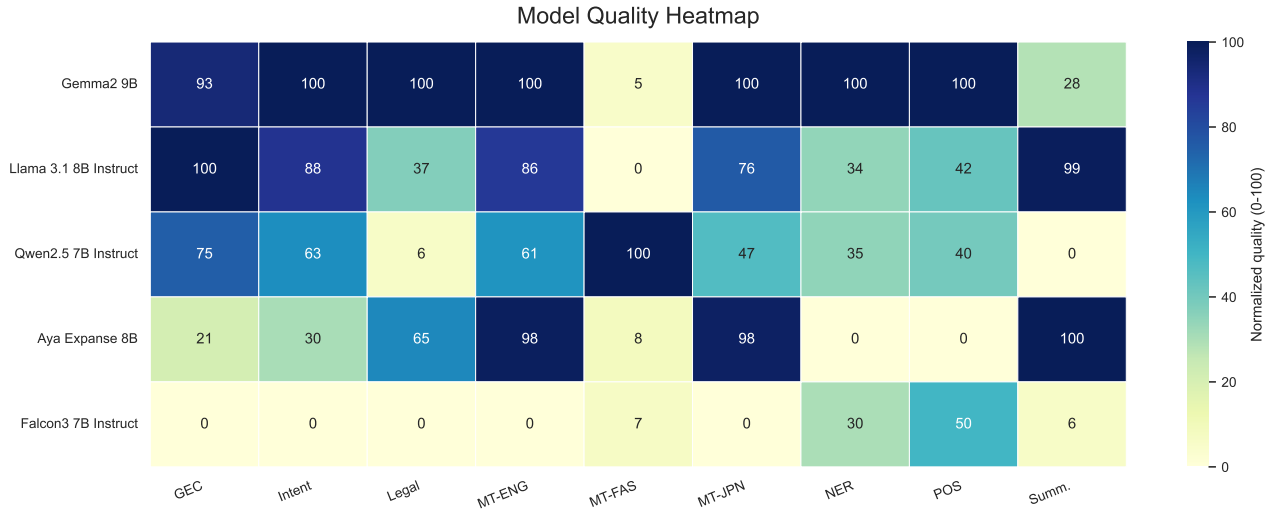

Figure 5: Normalized quality heatmap across the nine benchmark segments. Scores are normalized within each segment so that higher is always better, even when the underlying metric is lower-is-better such as WER.

detailed in Table 8, using chrF as the primary metric, `gemma2:9b` leads English and Japanese translation, while `qwen2.5:7b-instruct` is first on Persian. Notably, its chrF score substantially exceeds the rest of the field despite weaker WER, highlighting differences between surface-form and character-level evaluation metrics.

5.2.3 POS tagging and NER. As shown in Table 4 and the detailed results in Tables 10 and 9, `gemma2:9b` leads on macro-F1 in both POS tagging and NER. However, `llama3.1:8b-instruct` achieves higher Accuracy in NER (but not in POS tagging), suggesting differences in handling minority entity classes and class imbalance.

5.2.4 Summarization. The winner is `aya-expanse:8b`, but only by a very small margin from `llama3.1:8b-instruct` (0.519 vs. 0.517), which is also better in WER. This suggests that the top two models are effectively near-tied semantically.

6 DISCUSSION

The benchmark combines multiple Greek NLP settings in a single reproducible workflow, but the current evaluation still has limitations. The capped protocol is intentionally pragmatic rather than exhaustive. Although well suited for iterative comparison, it is not a substitute for full test-set evaluation when final archival numbers are required. Specific metric choices are also provisional. In particular, summarization and machine translation would benefit from stronger semantic metrics and human assessment, while legal classification would benefit from richer label semantics. Even so, the supplemental metrics improve interpretability relative to a WER-only or macro-F1-only summary.

Despite these caveats, the benchmark exposes meaningful structure in model behaviour. Some models offer strong general performance across multiple tasks, while others are better understood as specialists for translation or summarization. This kind of cross-task perspective is hard to obtain from isolated task runs and is exactly where a unified benchmark becomes valuable.

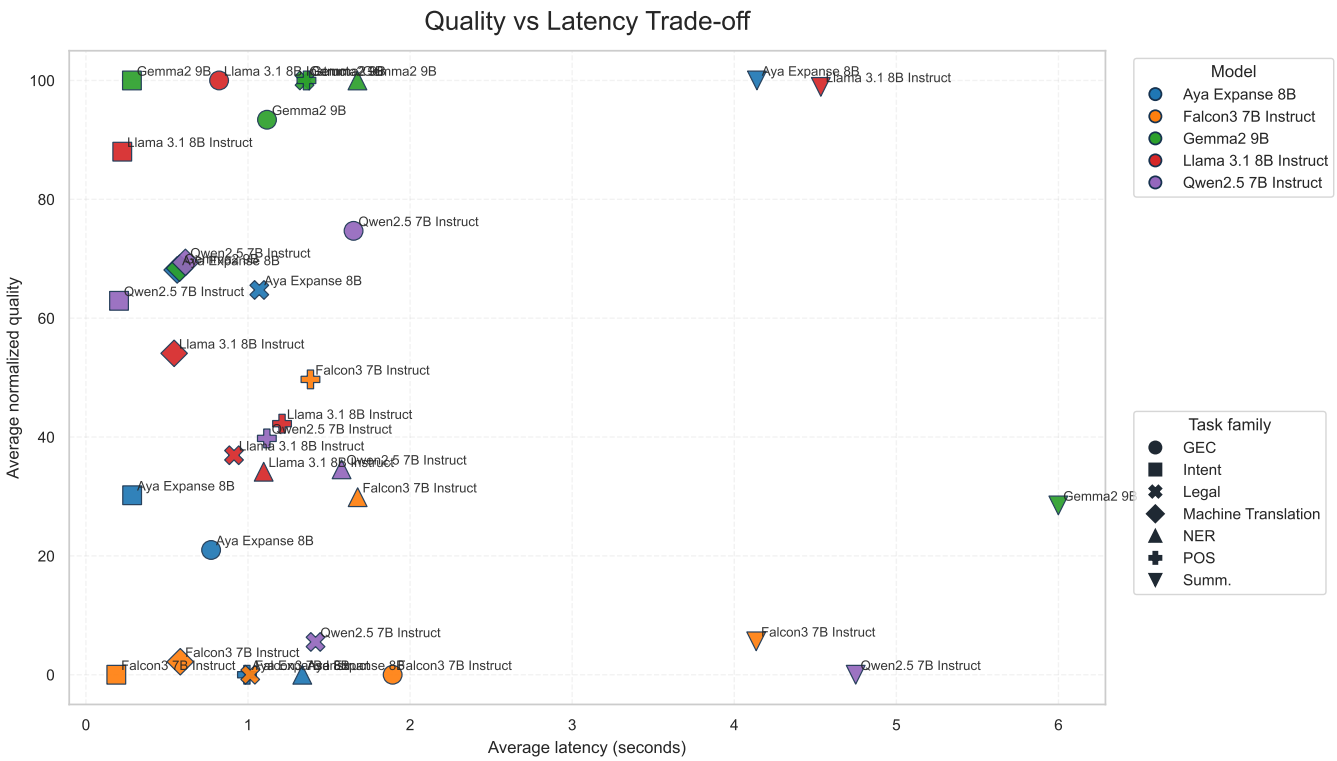

Figure 6: Quality-latency trade-off across models and task families. Each point corresponds to one model-task pair after averaging within the task family.

7 EVERGREEN DATASET DISCOVERY

Beyond our static benchmark, Greek NLP resources continue to evolve. To make the resource collection maintainable rather than fixed to the 2012–2023 survey, we implement a lightweight evergreen discovery pipeline for newly published work. This component is not used to define the benchmark scores reported above; instead, it demonstrates how the resource inventory can be updated as new datasets appear.

To support the continuous expansion of Greek NLP benchmarks, we adapt the search and screening protocol of ?] to recent literature (2024–2026) retrieved from the *ACL Anthology* and *Semantic Scholar* (n = 118).[15] The objective is to provide a proof of concept for an *evergreen* pipeline for identifying new Greek NLP datasets.

The screening process was applied in two stages. First, non–peer-reviewed publications were removed using rule-based filtering. The remaining exclusion criteria were then handled with a zero-shot LLM classifier. This process yielded 55 in-scope papers, of which 48 were identified as likely containing dataset references based on indications in the abstract that a dataset was introduced or released.

For these papers, dataset URLs were recovered automatically through a retrieval pipeline: article text was extracted, URLs were identified with regular expressions, and *Llama 3 70B* selected the most plausible dataset link. This produced 18 candidate dataset URLs, 13 of which were verified as reachable.

A manual audit on a small dataset-annotated subset (n = 18) showed high precision and moderate recall (precision = 1.00, recall = 0.69), indicating a conservative but reliable approach. Finally, titles and abstracts were classified into `canonical_task` categories using *Llama 3 70B*, enabling task-level organization of newly identified resources, mainly in text classification and machine translation.

Overall, the results indicate that automatic mining can effectively complement manual curation and support the systematic, continuous enrichment of Greek NLP resources.

8 CONCLUSION

We presented a reproducible Greek NLP benchmark for open, instruction-tuned, edge-capable LLMs and used it to compare five models across seven task families and nine benchmark segments. Under a deterministic capped-test protocol, `gemma2:9b` ranks first overall, while `llama3.1:8b-instruct`, `qwen2.5:7b-instruct`, and `aya-expanse:8b` emerge as specialists on GEC, Persian translation, and summarization respectively. The current pipeline is already useful for comparative evaluation, and it also provides a clean path from benchmark execution to publication assets. Future work should extend the benchmark with richer label semantics for legal classification, stronger semantic metrics and human assessment for translation

[15]The implementation is available in our GitHub repository, enabling reuse and continuous application of the protocol for future dataset discovery.

and summarization, and continuous dataset discovery to keep the Greek NLP resource collection up to date.

REFERENCES

[1] Georgios Alexandridis, Iraklis Varlamis, Konstantinos Korovesis, George Caridakis, and Panagiotis Tsantilas. 2021. A Survey on Sentiment Analysis and Opinion Mining in Greek Social Media. *Information* 12, 8 (2021). https://doi.org/10.3390/info12080331

[2] Ebtesam Almazrouei, Hamza Alobeidli, Abdulaziz Alshamsi, Alessandro Cappelli, Ruxandra Cojocaru, Mérouane Debbah, Étienne Goffinet, Daniel Hesslow, Julien Launay, Quentin Malartic, et al. 2023. The falcon series of open language models. *arXiv preprint arXiv:2311.16867* (2023).

[3] Nikos Bartziokas, Thanassis Mavropoulos, and Constantine Kotropoulos. 2020. Datasets and Performance Metrics for Greek Named Entity Recognition. In *11th Hellenic Conference on Artificial Intelligence* (Athens, Greece) *(SETN 2020)*. Association for Computing Machinery, New York, NY, USA, 160–167. https://doi.org/10.1145/3411408.3411437

[4] Vasileios Barzokas, Eirini Papagiannopoulou, and Grigorios Tsoumakas. 2020. Studying the Evolution of Greek Words via Word Embeddings. In *11th Hellenic Conference on Artificial Intelligence* (Athens, Greece) *(SETN 2020)*. Association for Computing Machinery, New York, NY, USA, 118–124. https://doi.org/10.1145/3411408.3411425

[5] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems* 33 (2020), 1877–1901.

[6] Ilias Chalkidis, Abhik Jana, Dirk Hartung, Michael Bommarito, Ion Androutsopoulos, Daniel Katz, and Nikolaos Aletras. 2022. LexGLUE: A Benchmark Dataset for Legal Language Understanding in English. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 4310–4330. https://doi.org/10.18653/v1/2022.acl-long.297

[7] John Dang, Shivalika Singh, Daniel D'souza, Arash Ahmadian, Alejandro Salamanca, Madeline Smith, Aidan Peppin, Sungjin Hong, Manoj Govindassamy, Terrence Zhao, et al. 2024. Aya expanse: Combining research breakthroughs for a new multilingual frontier. *arXiv preprint arXiv:2412.04261* (2024).

[8] Konstantina Dritsa, Aikaterini Thoma, Ioannis Pavlopoulos, and Panos Louridas. 2022. A Greek Parliament Proceedings Dataset for Computational Linguistics and Political Analysis. *Advances in Neural Information Processing Systems* 35 (2022), 28874–28888.

[9] Iakovos Evdaimon, Hadi Abdine, Christos Xypolopoulos, Stamatis Outsios, Michalis Vazirgiannis, and Giorgos Stamou. 2023. GreekBART: The First Pretrained Greek Sequence-to-Sequence Model. *arXiv preprint arXiv:2304.00869* (2023).

[10] Fotios Fitsilis and George Mikros. 2021. Development and validation of a corpus of written parliamentary questions in the hellenic parliament. *Journal of Open Humanities Data* 7 (2021).

[11] Panagiotis Gakis, Christos Panagiotakopoulos, Kyriakos Sgarbas, Christos Tsalidis, and Vassilios Verykios. 2016. Design and construction of the Greek grammar checker. *Digital Scholarship in the Humanities* 32, 3 (07 2016), 554–576. https://doi.org/10.1093/llc/fqw025 arXiv:https://academic.oup.com/dsh/article-pdf/32/3/554/19646879/fqw025.pdf

[12] Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sidney Black, Guy Blasi, Anthony Casey, Nate Cheeseman, Juan Ciro, Sasha Crabtree, Nicholas Currier, Xingyu Dang, Clementine Decoopman, Chris Foster, Raphael Gontijo Lopes, Huan He, Kevin Hu, Brandon Knight, Jorge Larios, Logan Li, Xander Lin, Shawn Miller, Niklas Muennighoff, Aryan Naeini, Hyunwoo Nam, Sam Pan, Louis Schmid, Hailey Schoelkopf, Xiang Wang, Albert Webson, Mitchell Wortsman, Yuxi Yao, Stella Biderman, Douwe Kiela, Teven Le Scao, Colin Raffel, Andy Zhou, Dan Hendrycks, et al. 2023. A Framework for Few-Shot Language Model Evaluation. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*.

[13] Maria Giatsoglou, Manolis G. Vozalis, Konstantinos Diamantaras, Athena Vakali, George Sarigiannidis, and Konstantinos Ch. Chatzisavvas. 2017. Sentiment analysis leveraging emotions and word embeddings. *Expert Systems with Applications* 69 (2017), 214–224. https://doi.org/10.1016/j.eswa.2016.10.043

[14] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783* (2024).

[15] Patrick Juola and Efstathios Stamatatos. 2013. Overview of the Author Identification Task at PAN 2013. *CLEF (Working Notes)* 1179 (2013).

[16] Marios Koniaris, Dimitris Galanis, Eugenia Giannini, and Panayiotis Tsanakas. 2023. Evaluation of automatic legal text summarization techniques for greek case law. *Information* 14, 4 (2023), 250.

[17] Moshe Koppel and Noam Ordan. 2011. Translationese and its dialects. In *Proceedings of the 49th annual meeting of the association for computational linguistics: Human language technologies*. 1318–1326.

[18] Katerina Korre, Marita Chatzipanagiotou, and John Pavlopoulos. 2021. ELERRANT: Automatic grammatical error type classification for Greek. In *Proceedings of the International Conference on Recent Advances in Natural Language Processing (RANLP 2021)*. 708–717.

[19] John Koutsikakis, Ilias Chalkidis, Prodromos Malakasiotis, and Ion Androutsopoulos. 2020. GREEK-BERT: The Greeks Visiting Sesame Street. In *11th Hellenic Conference on Artificial Intelligence* (Athens, Greece) *(SETN 2020)*. Association for Computing Machinery, New York, NY, USA, 110–117. https://doi.org/10.1145/3411408.3411440

[20] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bonnie Yan, Ce Zhang, M. Cosgrove, Christopher D. Manning, Christopher Ré, and Tatsunori B. Hashimoto. 2023. Holistic Evaluation of Language Models. *Transactions on Machine Learning Research* (2023).

[21] Chin-Yew Lin. 2004. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*. 74–81.

[22] Martin Majliš and Zdeněk Žabokrtskỳ. 2012. Language richness of the web. In *Proceedings of the Eighth International Conference on Language Resources and Evaluation (LREC'12)*. 2927–2934.

[23] Courtney Napoles, Keisuke Sakaguchi, Matt Post, and Joel Tetreault. 2015. Ground truth for grammatical error correction metrics. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*. 588–593.

[24] Maria Nefeli Nikiforos, Yorghos Voutos, Anthi Drougani, Phivos Mylonas, and Katia Lida Kermanidis. 2021. The Modern Greek Language on the Social Web: A Survey of Data Sets and Mining Applications. *Data* 6 (2021), 52.

[25] Joakim Nivre, Marie-Catherine De Marneffe, Filip Ginter, Yoav Goldberg, Jan Hajic, Christopher D Manning, Ryan McDonald, Slav Petrov, Sampo Pyysalo, Natalia Silveira, et al. 2016. Universal dependencies v1: A multilingual treebank collection. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*. 1659–1666.

[26] Christos Papaloukas, Ilias Chalkidis, Konstantinos Athinaios, Despina Pantazi, and Manolis Koubarakis. 2021. Multi-granular Legal Topic Classification on Greek Legislation. In *Proceedings of the Natural Legal Language Processing Workshop 2021*. Association for Computational Linguistics, Punta Cana, Dominican Republic, 63–75. https://doi.org/10.18653/v1/2021.nllp-1.6

[27] Katerina Papantoniou, Vasilis Efthymiou, and Dimitris Plexousakis. 2023. Automating Benchmark Generation for Named Entity Recognition and Entity Linking. In *European Semantic Web Conference*. Springer, 143–148.

[28] John Pavlopoulos and Maria Konstantinidou. 2023. Computational authorship analysis of the Homeric poems. *International Journal of Digital Humanities* 5, 1 (2023), 45–64.

[29] Maja Popović. 2015. chrF: character n-gram F-score for automatic MT evaluation. In *Proceedings of the tenth workshop on statistical machine translation*. 392–395.

[30] Sameer S Pradhan, Eduard Hovy, Mitch Marcus, Martha Palmer, Lance Ramshaw, and Ralph Weischedel. 2007. Ontonotes: A unified relational semantic representation. In *International Conference on Semantic Computing (ICSC 2007)*. IEEE, 517–526.

[31] Prokopis Prokopidis and Harris Papageorgiou. 2017. Universal dependencies for Greek. In *Proceedings of the nodalida 2017 workshop on universal dependencies (udw 2017)*. 102–106.

[32] Prokopis Prokopidis, Vassilis Papavassiliou, and Stelios Piperidis. 2016. Parallel Global Voices: a Collection of Multilingual Corpora with Citizen Media Stories. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*. European Language Resources Association (ELRA), Portorož, Slovenia, 900–905. https://aclanthology.org/L16-1144

[33] Prokopis Prokopidis and Stelios Piperidis. 2020. A Neural NLP toolkit for Greek. In *11th Hellenic Conference on Artificial Intelligence*. 125–128.

[34] Sofia Rizou, Angelos Theofilatos, Antonia Paflioti, Eleni Pissari, Iraklis Varlamis, George Sarigiannidis, and K Ch Chatzisavvas. 2023. Efficient intent classification and entity recognition for university administrative services employing deep learning models. *Intelligent Systems with Applications* 19 (2023), 200247.

[35] Erik F Sang and Fien De Meulder. 2003. Introduction to the CoNLL-2003 shared task: Language-independent named entity recognition. *arXiv preprint cs/0306050* (2003).

[36] Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, et al. 2024. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118* (2024).

[37] Gorka Urbizu, Iñaki San Vicente, Xabier Saralegi, Rodrigo Agerri, and Aitor Soroa. 2022. BasqueGLUE: A natural language understanding benchmark for Basque. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference*. 1603–1612.

[38] Leon Voukoutis, Dimitris Roussis, Georgios Paraskevopoulos, Sokratis Sofianopoulos, Prokopis Prokopidis, Vassilis Papavasileiou, Athanasios Katsamanis, Stelios Piperidis, and Vassilis Katsouros. 2024. Meltemi: The first open Large Language Model for Greek. *arXiv preprint arXiv:2407.20743* (2024).

[39] Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. 2019. Superglue: A stickier benchmark for general-purpose language understanding systems. *Advances in neural information processing systems* 32 (2019).

[40] Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R Bowman. 2019. Glue: A multi-task benchmark and analysis platform for natural language understanding. In *7th International Conference on Learning Representations, ICLR 2019*.

[41] Bryan Wilie, Karissa Vincentio, Genta Indra Winata, Samuel Cahyawijaya, Xiaohong Li, Zhi Yuan Lim, Sidik Soleman, Rahmad Mahendra, Pascale Fung, Syafri Bahar, and Ayu Purwarianti. 2020. IndoNLU: Benchmark and Resources for Evaluating Indonesian Natural Language Understanding. In *Proceedings of the 1st Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 10th International Joint Conference on Natural Language Processing*, Kam-Fai Wong, Kevin Knight, and Hua Wu (Eds.). Association for Computational Linguistics, Suzhou, China, 843–857. https://aclanthology.org/2020.aacl-main.85

[42] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. 2024. Qwen2.5 Technical Report. *arXiv preprint arXiv:2412.15115* (2024). https://doi.org/10.48550/arXiv.2412.15115

[43] Marcos Zampieri, Preslav Nakov, Sara Rosenthal, Pepa Atanasova, Georgi Karadzhov, Hamdy Mubarak, Leon Derczynski, Zeses Pitenis, and Çagri Çöltekin. 2020. SemEval-2020 Task 12: Multilingual Offensive Language Identification in Social Media (OffensEval 2020). *CoRR* abs/2006.07235 (2020). arXiv:2006.07235 https://arxiv.org/abs/2006.07235

[44] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q Weinberger, and Yoav Artzi. 2019. Bertscore: Evaluating text generation with bert. *arXiv preprint arXiv:1904.09675* (2019).

A WINNER MARGINS BY SEGMENT

Figure 7 presents the gap between the best and second-best model for each benchmark segment. The largest winner margins occur on NER and POS tagging.

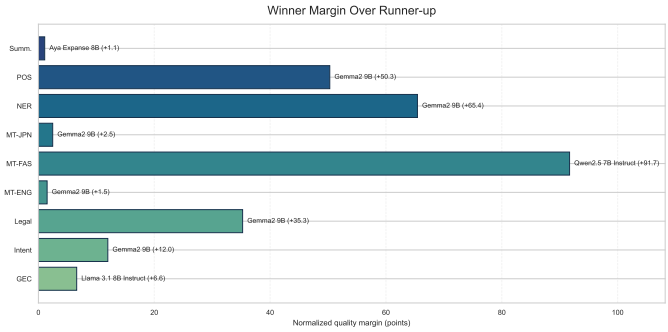

Figure 7: Normalized quality margin between the winning model and the runner-up for each benchmark segment. Large margins indicate clearer task-level separation between the top two models.

B PER-TASK ANALYSIS

B.1 GEC

B.2 Intent Classification

B.3 Legal Classification

B.4 Machine Translation

B.5 NER

B.6 POS

B.7 Summarization

C PROMPT TEMPLATES

The following tables list the system and user prompt templates used in the benchmark.

Table 5: Per-model results for GEC.

Model	Exact Match	WER	Latency (s)
Llama 3.1 8B Instruct	0.120	10.489	0.822
Gemma2 9B	0.137	11.670	1.117
Qwen2.5 7B Instruct	0.069	13.324	1.651
Aya Expanse 8B	0.057	21.025	0.772
Falcon3 7B Instruct	0.006	27.049	1.892

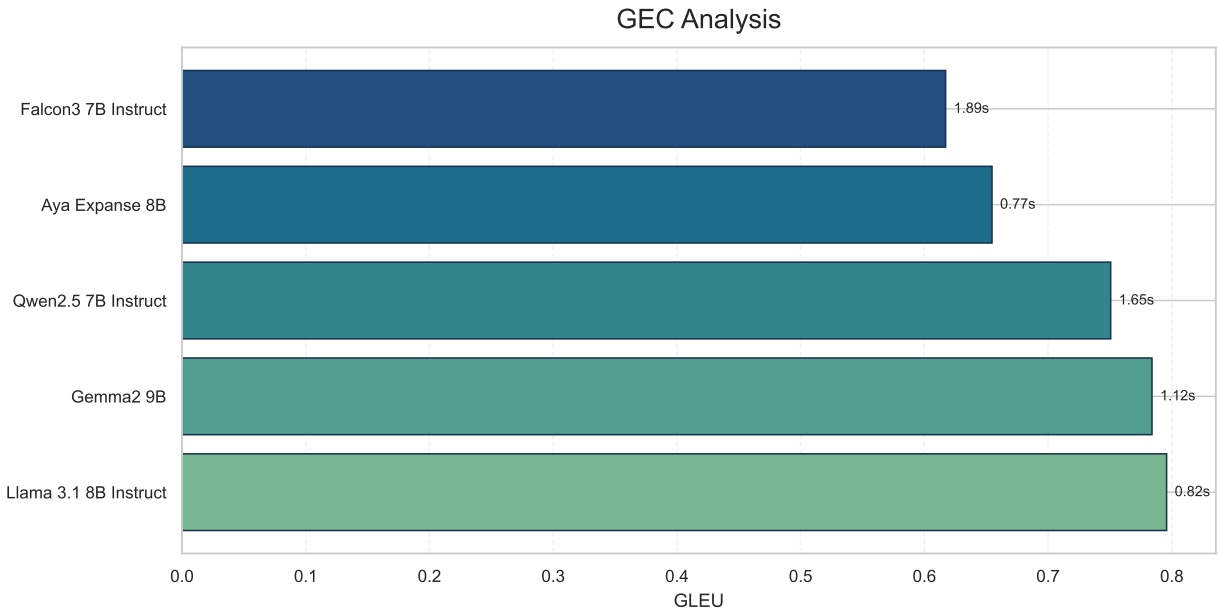

Figure 8: Per-task comparison for GEC. Bars show the primary metric and annotations show average latency.

Table 6: Per-model results for Intent.

Model	Accuracy	Macro-F1	Latency (s)
Gemma2 9B	0.862	0.737	0.284
Llama 3.1 8B Instruct	0.720	0.707	0.224
Qwen2.5 7B Instruct	0.752	0.645	0.204
Aya Expanse 8B	0.716	0.563	0.285
Falcon3 7B Instruct	0.578	0.488	0.188

Table 7: Per-model results for Legal.

Model	Accuracy	Macro-F1	Latency (s)
Gemma2 9B	0.346	0.126	1.349
Aya Expanse 8B	0.292	0.098	1.070
Llama 3.1 8B Instruct	0.182	0.076	0.914
Qwen2.5 7B Instruct	0.202	0.051	1.417
Falcon3 7B Instruct	0.122	0.047	1.014

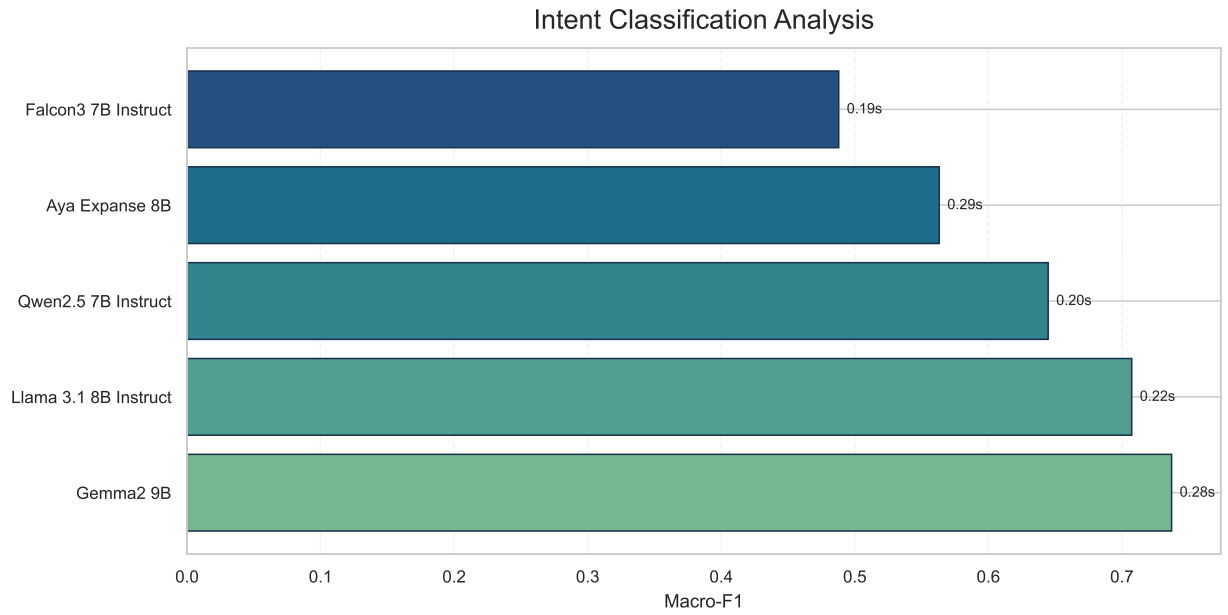

Figure 9: Per-task comparison for Intent Classification. Bars show the primary metric and annotations show average latency.

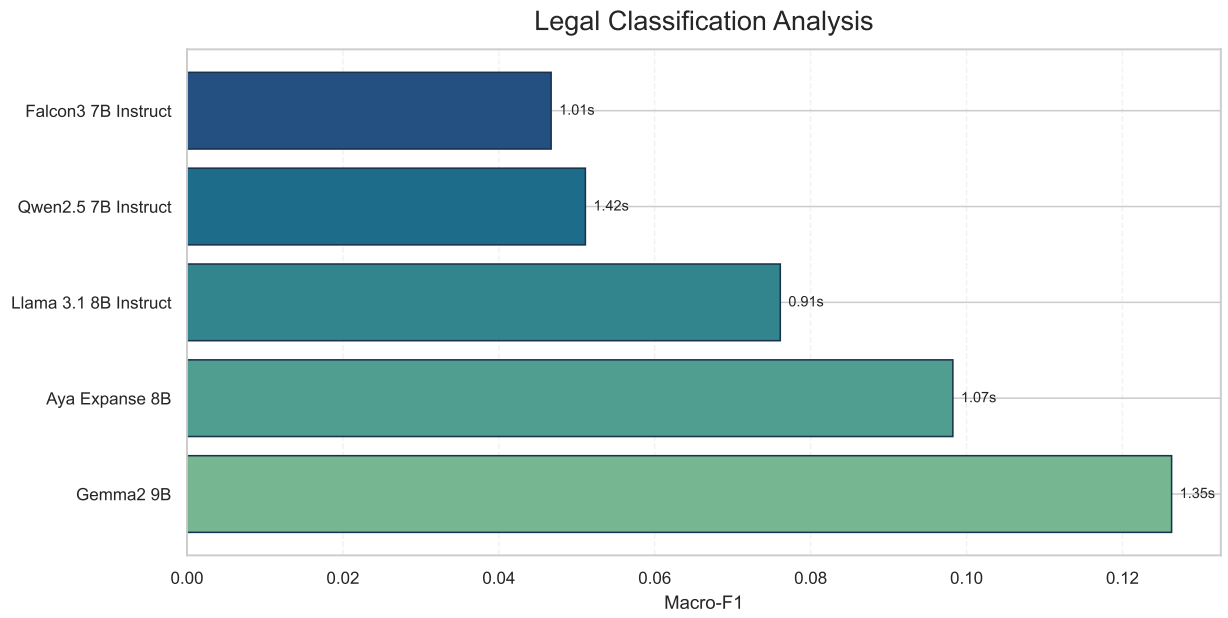

Figure 10: Per-task comparison for Legal Classification. Bars show the primary metric and annotations show average latency.

Table 8: Per-model results for Machine Translation.

Lang	Model	WER	BLEU	chrF	Latency (s)
ENG	Gemma2 9B	57.915	35.073	61.781	0.569
ENG	Aya Expanse 8B	56.359	35.037	61.591	0.531
ENG	Llama 3.1 8B Instruct	64.237	32.576	60.045	0.441
ENG	Qwen2.5 7B Instruct	66.484	29.280	56.907	0.439
ENG	Falcon3 7B Instruct	77.902	21.355	49.132	0.431
FAS	Qwen2.5 7B Instruct	116.337	4.183	18.325	0.845
JPN	Gemma2 9B	109.538	6.227	16.726	0.629
JPN	Aya Expanse 8B	109.067	2.279	16.533	0.610
JPN	Llama 3.1 8B Instruct	168.112	2.358	14.845	0.588
JPN	Qwen2.5 7B Instruct	168.901	3.656	12.553	0.558
JPN	Falcon3 7B Instruct	115.930	5.665	8.895	0.900
FAS	Aya Expanse 8B	109.222	2.146	3.408	0.548
FAS	Falcon3 7B Instruct	114.451	1.984	3.124	0.416
FAS	Gemma2 9B	110.071	1.926	2.906	0.536
FAS	Llama 3.1 8B Instruct	343.614	0.555	2.058	0.604

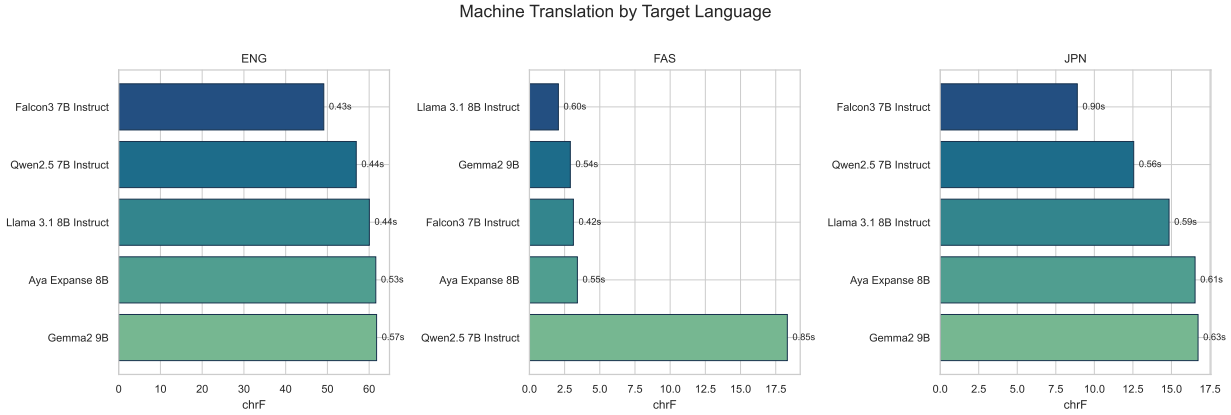

Figure 11: Per-task comparison for Machine Translation. Bars show the primary metric and annotations show average latency.

Table 9: Per-model results for NER.

Model	Accuracy	Macro-F1	Entity F1	Latency (s)
Gemma2 9B	0.796	0.122	0.043	1.675
Qwen2.5 7B Instruct	0.689	0.095	0.045	1.577
Llama 3.1 8B Instruct	0.812	0.094	0.049	1.097
Falcon3 7B Instruct	0.804	0.093	0.026	1.677
Aya Expanse 8B	0.735	0.080	0.010	1.335

Table 10: Per-model results for POS.

Model	Accuracy	Macro-F1	Latency (s)
Gemma2 9B	0.603	0.486	1.361
Falcon3 7B Instruct	0.438	0.332	1.385
Llama 3.1 8B Instruct	0.390	0.309	1.210
Qwen2.5 7B Instruct	0.354	0.302	1.116
Aya Expanse 8B	0.229	0.180	0.993

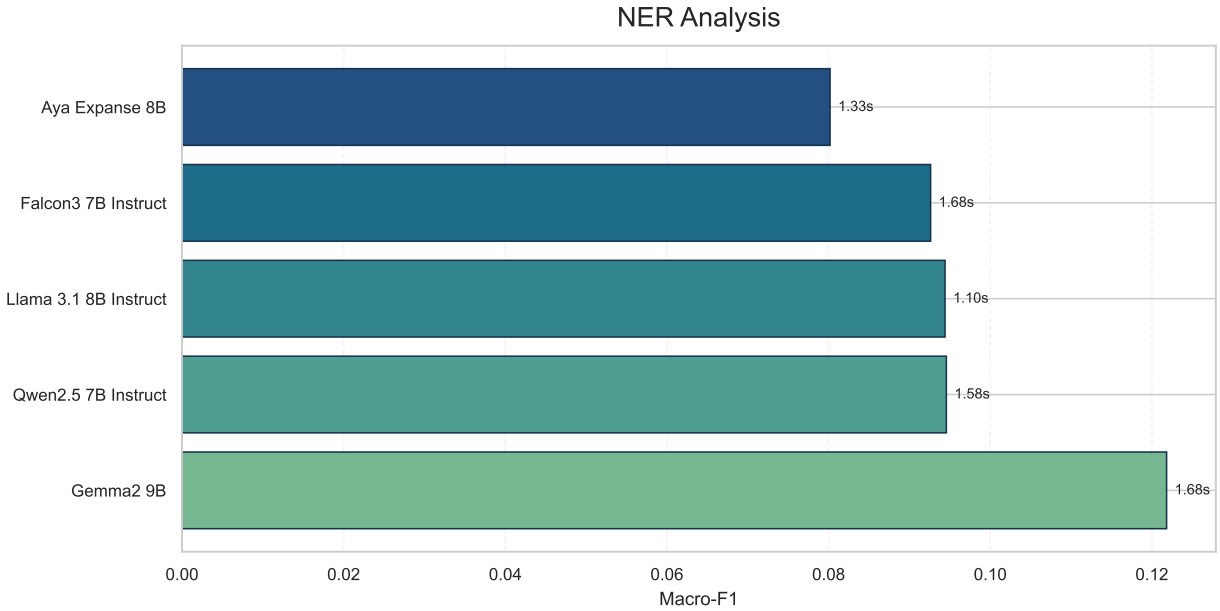

Figure 12: Per-task comparison for NER. Bars show the primary metric and annotations show average latency.

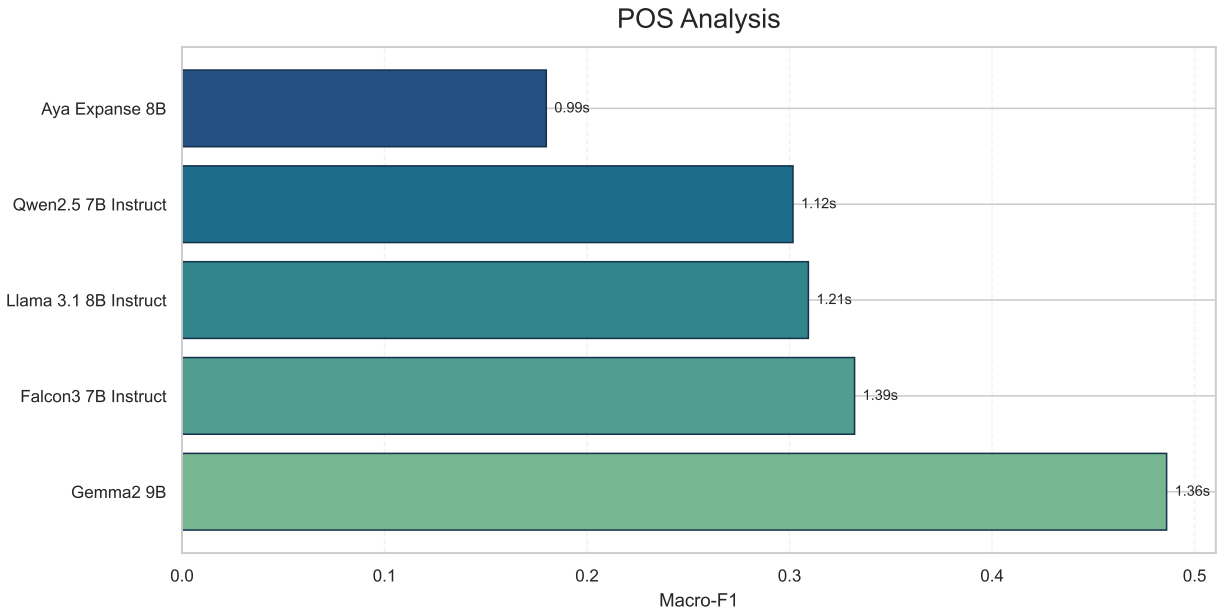

Figure 13: Per-task comparison for POS. Bars show the primary metric and annotations show average latency.

Table 11: Per-model results for Summ..

Model	WER	ROUGE-1	ROUGE-L	BERTScore-F1	Latency (s)
Aya Expanse 8B	208.784	0.064	0.063	0.519	4.140
Llama 3.1 8B Instruct	173.585	0.081	0.076	0.517	4.534
Gemma2 9B	271.162	0.028	0.027	0.407	5.999
Falcon3 7B Instruct	235.214	0.009	0.009	0.371	4.135
Qwen2.5 7B Instruct	327.113	0.016	0.014	0.362	4.749

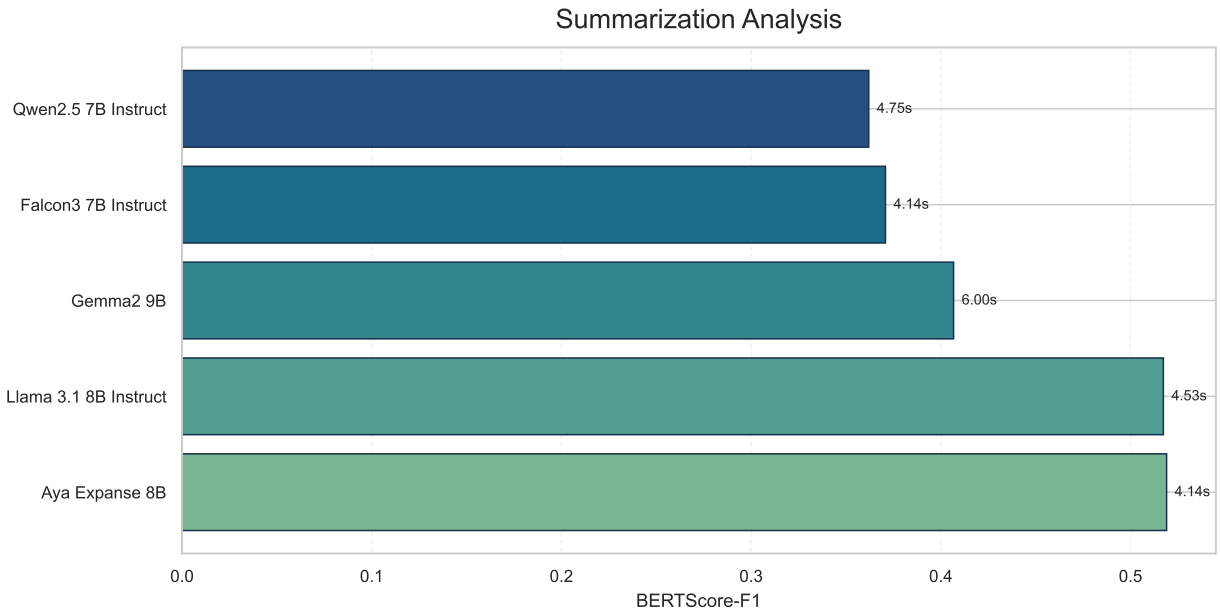

Figure 14: Per-task comparison for Summarization. Bars show the primary metric and annotations show average latency.

Table 17: POS Tagging.

System	You are a Greek POS tagging system. Return only a Python-style list of tags and nothing else.
User	You are given a tokenized Greek sentence. Return only a Python-style list of POS tags, one tag per token. Use only these labels: <POS_LABEL_LIST> Tokens: <TOKEN_LIST> Tags:

Table 18: Summarization.

System	You are a Greek summarization system. Produce a concise summary in Greek and return only the summary.
User	Given a Greek legal text, provide its summary also in Greek. Only the summary should be returned. Text: <DOCUMENT_TEXT> Summary: